

Looking at your SignConnect project, I'll provide a comprehensive senior-level review focusing on architecture, best practices, and areas for improvement.

Overall Assessment

You've built a solid full-stack application with good fundamentals - Docker containerization, proper testing setup, clean separation of concerns, and modern tech choices. The project shows strong intermediate-level engineering skills. Let's dive into specific areas where you can elevate this to senior-level work.

What You're Doing Well

1. **Modern Tech Stack:** FastAPI + React + PostgreSQL with pgvector is excellent for your use case
2. **Testing Infrastructure:** Good pytest setup with fixtures and Docker integration
3. **Authentication:** Smart choice using Firebase for auth instead of rolling your own
4. **Vector Search:** Using pgvector for semantic search shows forward-thinking
5. **Structured Logging:** Using structlog is a professional touch
6. **Database Migrations:** Alembic setup shows maturity

Critical Issues to Address

1. Security Vulnerabilities

python

```
# In websockets.py - Token is exposed in URL
@router.websocket("/ws")
async def websocket_endpoint(
    websocket: WebSocket,
    db: Session = Depends(get_db),
    user: dict = Depends(get_current_user_ws), # Token in query param!
):
```

Problem: Tokens in URLs are logged by proxies/servers and visible in browser history.

Solution: Send token as first message after connection:

python

```
@router.websocket("/ws")
async def websocket_endpoint(websocket: WebSocket, db: Session = Depends(get_db)):
    await websocket.accept()
    try:
        # First message should be authentication
        auth_message = await websocket.receive_json()
        if auth_message.get("type") != "auth":
            await websocket.close(code=1008, reason="Authentication required")
            return

        token = auth_message.get("token")
        user = verify_firebase_token(token)
        if not user:
            await websocket.close(code=1008, reason="Invalid token")
            return

        # Now proceed with authenticated connection
        await handle_authenticated_connection(websocket, user, db)
    except WebSocketDisconnect:
        pass
```

2. Missing Production Configurations

Your `.env` files expose sensitive defaults. Create proper environment-specific configs:

python

```
# src/signconnect/core/config.py
from enum import Enum

class Environment(str, Enum):
    DEVELOPMENT = "development"
    STAGING = "staging"
    PRODUCTION = "production"

class Settings(BaseSettings):
    ENVIRONMENT: Environment = Environment.DEVELOPMENT

    # Add rate limiting
    RATE_LIMIT_ENABLED: bool = Field(True, env="RATE_LIMIT_ENABLED")
    RATE_LIMIT_REQUESTS: int = Field(100, env="RATE_LIMIT_REQUESTS")
    RATE_LIMIT_PERIOD: int = Field(60, env="RATE_LIMIT_PERIOD") # seconds

    # Add CORS configuration
    CORS_ORIGINS: List[str] = Field(
        default_factory=lambda: ["http://localhost:5173"],
        env="CORS_ORIGINS"
    )

    @computed_field
    @property
    def is_production(self) -> bool:
        return self.ENVIRONMENT == Environment.PRODUCTION
```

3. Database Connection Pool Issues

You're creating engines without proper pooling configuration:

python

```
# Better approach in app_factory.py
from sqlalchemy.pool import NullPool, QueuePool

def create_engine_with_pool(settings: Settings):
    if settings.ENVIRONMENT == Environment.PRODUCTION:
        return create_engine(
            str(settings.DATABASE_URL),
            poolclass=QueuePool,
            pool_size=20,
            max_overflow=40,
            pool_timeout=30,
            pool_pre_ping=True, # Verify connections before using
        )
    else:
        return create_engine(
            str(settings.DATABASE_URL),
            poolclass=NullPool # No pooling for development
        )
```

Architecture Improvements

1. Implement Service Layer Pattern

You're mixing business logic with route handlers. Create a service layer:

python

```
# src/signconnect/services/scenario_service.py
from typing import List, Optional
import uuid
from sqlalchemy.orm import Session
from .. import crud, schemas
from ..exceptions import NotFoundError, ForbiddenError

class ScenarioService:
    def __init__(self, db: Session):
        self.db = db

    def create_scenario(
        self,
        user_id: uuid.UUID,
        scenario_data: schemas.ScenarioCreate
    ) -> schemas.Scenario:
        # Business logic here
        existing = crud.get_scenario_by_name(
            self.db,
            name=scenario_data.name,
            user_id=user_id
        )
        if existing:
            raise ValueError(f"Scenario '{scenario_data.name}' already exists")

        return crud.create_scenario(
            self.db,
            scenario=scenario_data,
            user_id=user_id
        )

    def get_user_scenarios(
        self,
        user_id: uuid.UUID
    ) -> List[schemas.Scenario]:
        return crud.get_scenarios_by_user(self.db, user_id)

# Then in routers/scenarios.py
@router.post("/", response_model=schemas.Scenario)
def create_scenario(
    scenario: schemas.ScenarioCreate,
    db: Session = Depends(get_db),
    current_user: dict = Depends(get_current_user),
    service: ScenarioService = Depends(get_scenario_service),
):
    return service.create_scenario(
        user_id=current_user.id,
        scenario_data=scenario
    )
```

2. Add Repository Pattern for Data Access

Replace CRUD functions with repositories:

python

```
# src/signconnect/repositories/scenario_repository.py
from typing import List, Optional
import uuid
from sqlalchemy.orm import Session
from ..db import models

class ScenarioRepository:
```

```

def __init__(self, db: Session):
    self.db = db

def find_by_id(self, scenario_id: uuid.UUID) -> Optional[models.Scenario]:
    return self.db.query(models.Scenario).filter(
        models.Scenario.id == scenario_id
    ).first()

def find_by_user(self, user_id: uuid.UUID) -> List[models.Scenario]:
    return self.db.query(models.Scenario).filter(
        models.Scenario.user_id == user_id
    ).all()

def save(self, scenario: models.Scenario) -> models.Scenario:
    self.db.add(scenario)
    self.db.commit()
    self.db.refresh(scenario)
    return scenario

def delete(self, scenario: models.Scenario) -> None:
    self.db.delete(scenario)
    self.db.commit()

```

3. Implement Proper Error Handling

Create custom exceptions and global handlers:

python

```

# src/signconnect/exceptions.py
class SignConnectException(Exception):
    """Base exception for all custom exceptions"""
    pass

class NotFoundError(SignConnectException):
    """Resource not found"""
    pass

class ForbiddenError(SignConnectException):
    """Access forbidden"""
    pass

class ConflictError(SignConnectException):
    """Resource conflict"""
    pass

# In app_factory.py
from fastapi import Request
from fastapi.responses import JSONResponse

@app.exception_handler(NotFoundError)
async def not_found_handler(request: Request, exc: NotFoundError):
    return JSONResponse(
        status_code=404,
        content={"detail": str(exc)}
    )

@app.exception_handler(ConflictError)
async def conflict_handler(request: Request, exc: ConflictError):
    return JSONResponse(
        status_code=409,
        content={"detail": str(exc)}
    )

```

Performance & Scalability Issues

1. WebSocket Connection Management

Your current implementation lacks connection pooling and cleanup:

python

```
# src/signconnect/services/connection_manager.py
from typing import Dict, Set
import asyncio
from datetime import datetime, timedelta

class EnhancedConnectionManager:
    def __init__(self):
        self.active_connections: Dict[str, WebSocket] = {}
        self.connection_metadata: Dict[str, dict] = {}
        self.cleanup_task = None

    async def connect(self, user_id: str, websocket: WebSocket):
        await websocket.accept()

        # Disconnect existing connection for this user
        if user_id in self.active_connections:
            await self.disconnect(user_id, reason="New connection established")

        self.active_connections[user_id] = websocket
        self.connection_metadata[user_id] = {
            "connected_at": datetime.utcnow(),
            "last_activity": datetime.utcnow()
        }

        # Start cleanup task if not running
        if not self.cleanup_task:
            self.cleanup_task = asyncio.create_task(self._cleanup_stale_connections())

    async def _cleanup_stale_connections(self):
        """Remove inactive connections every 30 seconds"""
        while self.active_connections:
            await asyncio.sleep(30)
            now = datetime.utcnow()
            stale_threshold = now - timedelta(minutes=5)

            for user_id in list(self.connection_metadata.keys()):
                metadata = self.connection_metadata.get(user_id)
                if metadata and metadata["last_activity"] < stale_threshold:
                    await self.disconnect(user_id, reason="Inactive connection")
```

2. Add Caching Layer

Implement Redis for caching frequently accessed data:

python

```
# src/signconnect/services/cache_service.py
import redis.asyncio as redis
import json
from typing import Optional, Any
from datetime import timedelta

class CacheService:
    def __init__(self, redis_url: str):
        self.redis = redis.from_url(redis_url)

    async def get(self, key: str) -> Optional[Any]:
        value = await self.redis.get(key)
        return json.loads(value) if value else None
```

```

    async def set(
        self,
        key: str,
        value: Any,
        ttl: Optional[timedelta] = None
    ):
        serialized = json.dumps(value)
        if ttl:
            await self.redis.setex(key, ttl, serialized)
        else:
            await self.redis.set(key, serialized)

    async def invalidate(self, pattern: str):
        """Invalidate all keys matching pattern"""
        keys = await self.redis.keys(pattern)
        if keys:
            await self.redis.delete(*keys)

# Usage in service layer
class ScenarioService:
    def __init__(self, db: Session, cache: CacheService):
        self.db = db
        self.cache = cache

    async def get_user_scenarios(self, user_id: uuid.UUID) -> List[schemas.Scenario]:
        cache_key = f"user:{user_id}:scenarios"

        # Try cache first
        cached = await self.cache.get(cache_key)
        if cached:
            return [schemas.Scenario(**s) for s in cached]

        # Fetch from DB
        scenarios = crud.get_scenarios_by_user(self.db, user_id)

        # Cache for 5 minutes
        await self.cache.set(
            cache_key,
            [s.dict() for s in scenarios],
            ttl=timedelta(minutes=5)
        )

        return scenarios

```

Production Readiness Gaps

1. Add API Rate Limiting

python

```

# src/signconnect/middleware/rate_limit.py
from fastapi import Request, HTTPException
from slowapi import Limiter
from slowapi.util import get_remote_address

limiter = Limiter(key_func=get_remote_address)

# In app_factory.py
from slowapi import _rate_limit_exceeded_handler
from slowapi.errors import RateLimitExceeded

app.state.limiter = limiter
app.add_exception_handler(RateLimitExceeded, _rate_limit_exceeded_handler)

# In routes

```

```
@router.post("/preferences/")
@limiter.limit("10/minute")
async def create_preference(request: Request, ...):
    pass
```

2. Add Health Checks

python

```
# src/signconnect/routers/health.py
from fastapi import APIRouter, Depends
from sqlalchemy import text
from sqlalchemy.orm import Session
from ..dependencies import get_db

router = APIRouter(tags=["health"])

@router.get("/health")
async def health_check():
    return {"status": "healthy"}

@router.get("/health/ready")
async def readiness_check(db: Session = Depends(get_db)):
    try:
        # Check database connectivity
        db.execute(text("SELECT 1"))

        # Check other critical services
        # ... check Redis, external APIs, etc.

        return {"status": "ready"}
    except Exception as e:
        return JSONResponse(
            status_code=503,
            content={"status": "not ready", "error": str(e)}
        )
```

3. Add Request ID Tracking

python

```
# src/signconnect/middleware/request_id.py
import uuid
from fastapi import Request
from starlette.middleware.base import BaseHTTPMiddleware

class RequestIDMiddleware(BaseHTTPMiddleware):
    async def dispatch(self, request: Request, call_next):
        request_id = str(uuid.uuid4())
        request.state.request_id = request_id

        response = await call_next(request)
        response.headers["X-Request-ID"] = request_id

        return response

# Add to structured logging
logger.bind(request_id=request.state.request_id)
```

Testing Improvements

1. Add Integration Tests for WebSockets

python

```
# tests/integration/test_websocket_flow.py
import pytest
from fastapi.testclient import TestClient
import json

@pytest.mark.asyncio
async def test_websocket_full_flow(authenticated_client: TestClient):
    with authenticated_client.websocket_connect("/api/ws") as websocket:
        # Send auth
        websocket.send_json({"type": "auth", "token": "fake-token"})

        # Send audio
        websocket.send_json({
            "type": "audio",
            "data": base64.b64encode(b"fake-audio").decode()
        })

        # Receive transcript
        data = websocket.receive_json()
        assert data["type"] in ["interim_transcript", "final_transcript"]

        # Request suggestions
        websocket.send_json({
            "type": "get_suggestions",
            "transcript": "test transcript"
        })

        # Receive suggestions
        data = websocket.receive_json()
        assert data["type"] == "suggestions"
        assert isinstance(data["data"], list)
```

2. Add Performance Tests

python

```
# tests/performance/test_load.py
import pytest
from locust import HttpUser, task, between

class SignConnectUser(HttpUser):
    wait_time = between(1, 3)

    def on_start(self):
        # Authenticate
        self.client.headers["Authorization"] = "Bearer fake-token"

    @task(3)
    def get_scenarios(self):
        self.client.get("/api/users/me/scenarios/")

    @task(1)
    def create_scenario(self):
        self.client.post("/api/users/me/scenarios/", json={
            "name": f"Test Scenario {uuid.uuid4()}",
            "description": "Load test scenario"
        })
```



Monitoring & Observability

1. Add Prometheus Metrics

python

```
# src/signconnect/metrics.py
from prometheus_client import Counter, Histogram, Gauge
import time

# Define metrics
http_requests_total = Counter(
    'http_requests_total',
    'Total HTTP requests',
    ['method', 'endpoint', 'status']
)

request_duration_seconds = Histogram(
    'request_duration_seconds',
    'HTTP request duration',
    ['method', 'endpoint']
)

active_websocket_connections = Gauge(
    'active_websocket_connections',
    'Number of active WebSocket connections'
)

# Middleware to track metrics
class MetricsMiddleware(BaseHTTPMiddleware):
    async def dispatch(self, request: Request, call_next):
        start_time = time.time()

        response = await call_next(request)

        duration = time.time() - start_time
        http_requests_total.labels(
            method=request.method,
            endpoint=request.url.path,
            status=response.status_code
        ).inc()

        request_duration_seconds.labels(
            method=request.method,
            endpoint=request.url.path
        ).observe(duration)

        return response
```

2. Add OpenTelemetry Tracing

python

```
# src/signconnect/tracing.py
from opentelemetry import trace
from opentelemetry.exporter.otlp.proto.grpc.trace_exporter import OTLPSpanExporter
from opentelemetry.sdk.trace import TracerProvider
from opentelemetry.sdk.trace.export import BatchSpanProcessor
from opentelemetry.instrumentation.fastapi import FastAPIInstrumentor
from opentelemetry.instrumentation.sqlalchemy import SQLAlchemyInstrumentor

def setup_tracing(app: FastAPI, settings: Settings):
    if not settings.is_production:
        return

    trace.set_tracer_provider(TracerProvider())
    tracer = trace.get_tracer(__name__)

    otlp_exporter = OTLPSpanExporter(
```

```

        endpoint=settings.OTLP_ENDPOINT,
        insecure=True
    )

    span_processor = BatchSpanProcessor(otlp_exporter)
    trace.get_tracer_provider().add_span_processor(span_processor)

    # Instrument FastAPI
    FastAPIInstrumentor.instrument_app(app)

    # Instrument SQLAlchemy
    SQLAlchemyInstrumentor().instrument(
        engine=engine,
        service="signconnect-db"
    )

```

Deployment Improvements

1. Create Proper Kubernetes Manifests

Since you have a K3s cluster, here's a production-ready deployment:

yaml

```

# k8s/signconnect/deployment.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: signconnect-api
  namespace: signconnect
spec:
  replicas: 3
  strategy:
    type: RollingUpdate
    rollingUpdate:
      maxSurge: 1
      maxUnavailable: 0
  selector:
    matchLabels:
      app: signconnect-api
  template:
    metadata:
      labels:
        app: signconnect-api
    spec:
      containers:
        - name: api
          image: ghcr.io/howdiggy/signconnect:latest
          ports:
            - containerPort: 8000
          env:
            - name: DATABASE_URL
              valueFrom:
                secretKeyRef:
                  name: signconnect-secrets
                  key: database-url
            - name: REDIS_URL
              valueFrom:
                secretKeyRef:
                  name: signconnect-secrets
                  key: redis-url
          resources:
            requests:
              memory: "256Mi"
              cpu: "100m"
            limits:

```

```

        memory: "512Mi"
        cpu: "500m"
    livenessProbe:
        httpGet:
            path: /health
            port: 8000
        initialDelaySeconds: 30
        periodSeconds: 10
    readinessProbe:
        httpGet:
            path: /health/ready
            port: 8000
        initialDelaySeconds: 5
        periodSeconds: 5
---
apiVersion: v1
kind: Service
metadata:
  name: signconnect-api
  namespace: signconnect
spec:
  selector:
    app: signconnect-api
  ports:
    - port: 80
      targetPort: 8000
  type: ClusterIP
---
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: signconnect-api
  namespace: signconnect
  annotations:
    cert-manager.io/cluster-issuer: letsencrypt-prod
    nginx.ingress.kubernetes.io/websocket-services: signconnect-api
    nginx.ingress.kubernetes.io/proxy-read-timeout: "3600"
    nginx.ingress.kubernetes.io/proxy-send-timeout: "3600"
spec:
  tls:
    - hosts:
        - api.signconnect.example.com
      secretName: signconnect-tls
  rules:
    - host: api.signconnect.example.com
      http:
        paths:
          - path: /
            pathType: Prefix
            backend:
              service:
                name: signconnect-api
                port:
                  number: 80

```

2. Add ArgoCD Application

yaml

```

# argocd/applications/signconnect.yaml
apiVersion: argoproj.io/v1alpha1
kind: Application
metadata:
  name: signconnect
  namespace: argocd
spec:

```

```

project: default
source:
  repoURL: https://github.com/HowDiggy/SignConnect
  targetRevision: main
  path: k8s/signconnect
destination:
  server: https://kubernetes.default.svc
  namespace: signconnect
syncPolicy:
  automated:
    prune: true
    selfHeal: true
  syncOptions:
    - CreateNamespace=true

```



Code Quality Improvements

1. Add Type Hints Everywhere

python

```

# Currently missing in many places
from typing import List, Optional, Dict, Any
from uuid import UUID

async def get_response_suggestions(
    transcript: str,
    user_preferences: List[str],
    conversation_history: List[str],
) -> List[str]: # Good!
    pass

# Add to all functions
def create_scenario(
    db: Session,
    scenario: schemas.ScenarioCreate,
    user_id: UUID
) -> models.Scenario:
    pass

```

2. Add Docstrings Following Google Style

python

```

def find_similar_question(
    db: Session,
    query_text: str,
    user_id: UUID
) -> Optional[models.ScenarioQuestion]:
    """Find the most similar question for a user using vector similarity.

    Args:
        db: Database session
        query_text: Text to search for similar questions
        user_id: ID of the user whose questions to search

    Returns:
        The most similar ScenarioQuestion or None if no questions exist

    Raises:
        DatabaseError: If the vector search operation fails
    """
    pass

```

3. Add Pre-commit Hooks

yaml

```
# .pre-commit-config.yaml
repos:
- repo: https://github.com/pre-commit/pre-commit-hooks
  rev: v4.6.0
  hooks:
  - id: check-yaml
  - id: end-of-file-fixer
  - id: trailing-whitespace
  - id: check-added-large-files

- repo: https://github.com/astral-sh/ruff-pre-commit
  rev: v0.4.8
  hooks:
  - id: ruff
    args: [--fix]
  - id: ruff-format

- repo: https://github.com/pre-commit/mirrors-mypy
  rev: v1.10.0
  hooks:
  - id: mypy
    additional_dependencies: [types-all]

- repo: https://github.com/sqlfluff/sqlfluff
  rev: 2.0.0
  hooks:
  - id: sqlfluff-lint
  - id: sqlfluff-fix
```

Immediate Action Items (Priority Order)

1. **Fix WebSocket security** - Move token from URL to message
2. **Add connection pooling** - Configure SQLAlchemy properly
3. **Implement rate limiting** - Protect your APIs
4. **Add health checks** - Critical for Kubernetes
5. **Set up Redis caching** - Improve performance
6. **Create service layer** - Separate business logic
7. **Add request tracing** - For debugging production issues
8. **Implement proper error handling** - Better user experience
9. **Add comprehensive logging** - You have structlog, use it more
10. **Write integration tests for WebSockets** - Critical path testing

Senior Engineer Tips

1. **Think in Terms of Failure:** Every external call can fail. Handle it gracefully.
2. **Design for Observability:** You can't fix what you can't see. Add metrics everywhere.
3. **Optimize for Change:** Use dependency injection, interfaces, and clear boundaries.
4. **Document Decisions:** Add ADRs (Architecture Decision Records) to your repo.
5. **Consider Data Privacy:** You're handling accessibility data - add data retention policies.

What Would Make This Portfolio-Ready

1. Add a comprehensive README with architecture diagrams
2. Include performance benchmarks (requests/second, latency percentiles)
3. Add a CI/CD pipeline with automated deployments
4. Create a demo environment with sample data
5. Add API documentation with OpenAPI/Swagger

6. Include monitoring dashboards (Grafana screenshots)
7. Write a blog post about your vector search implementation
8. Add load testing results
9. Include security scanning reports
10. Create a roadmap for future features

This is a solid project that demonstrates good fundamentals. With these improvements, you'll have a production-ready application that showcases senior-level engineering skills. Focus on the critical security and performance issues first, then gradually implement the architectural improvements. Remember, the goal is to show you can build systems that are not just functional, but reliable, scalable, and maintainable in a production environment.